



# Large Scale Data Processing

## Lecture 4 – Spark - streaming

dr hab. inż. Tomasz Kajdanowicz, Piotr Bielak, Roman  
Bartusiak, Krzysztof Rajda

November 18, 2021

Spark - streaming




# Architecture

## Spark - streaming



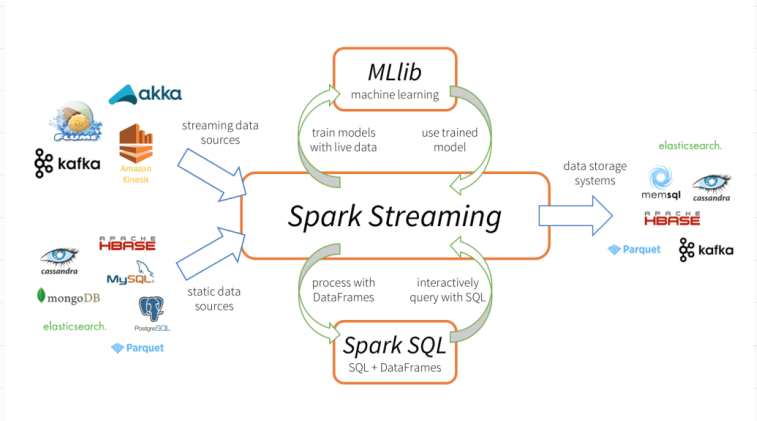
# Architecture

## Spark - streaming



# Architecture

## Spark - streaming



# Concepts - DStream

## Spark - streaming

- ▶ Utilizes low-level Spark API
- ▶ discretized-stream
- ▶ micro-batches
- ▶ "old"

# Concepts - Structured Streaming

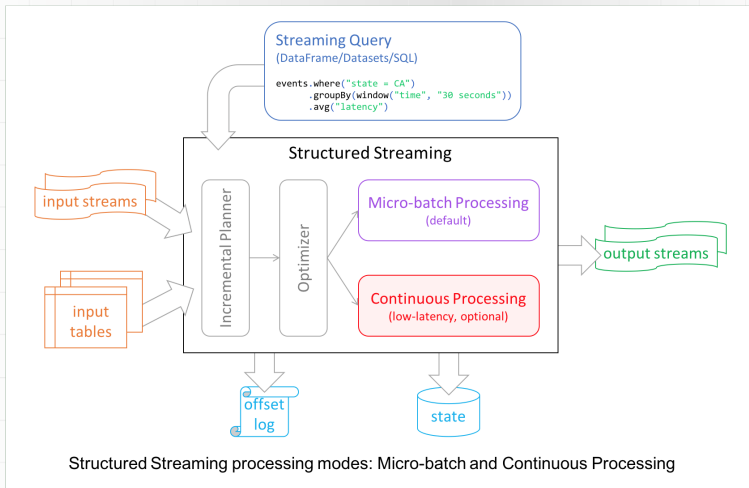
## Spark - streaming

- ▶ On top of Spark SQL
- ▶ batches and continuous
- ▶ streaming DataSets
- ▶ streaming DataFrames
- ▶ "new"




# Concepts - Structured Streaming

## Spark - streaming



# Concepts - Structured Streaming


## Spark - streaming



Micro-batch Processing uses periodic tasks to process events

# Concepts - Structured Streaming

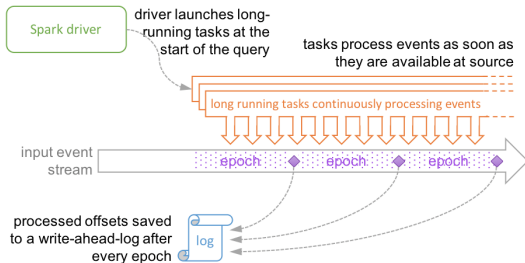
## Spark - streaming



Second-scale end-to-end latencies with Micro-batch Processing

# Concepts - Structured Streaming


## Spark - streaming



Continuous Processing uses long-running tasks to continuously process events

# Concepts - Structured Streaming

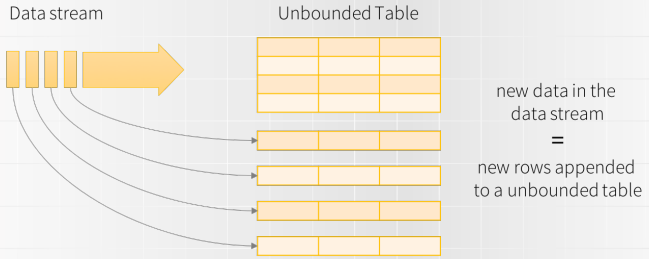
## Spark - streaming



Millisecond-scale end-to-end latencies with Continuous Processing

# Concepts - Data frames

## Spark - streaming



Data stream as an unbounded table

# Concepts - Input and Receiver

## Spark - streaming

- ▶ Each data input is consuming resources
- ▶ Each input is connected with receiver
- ▶ Receiver is responsible for fetching data and pushing to Spark
- ▶ Number of cores must be tuned to number of receivers
- ▶ Too low number of cores - no resources to process data as all resources will be occupied by receivers

# Concepts - Kafka

## Spark - streaming

- ▶ Distributed event streaming platform
- ▶ Events - records (key+value+timestamp)
- ▶ Events are organized in topics
- ▶ Producers and consumers
- ▶ Each topic is partitioned by event key - scalability
- ▶ Events in partitions are guaranteed to be returned to consumers in write order
- ▶ Offset - tracking already consumed events by consumer in given partition
- ▶ Data can be replicated



# Concepts - Kafka

## Spark - streaming

- ▶ Kafka Clients need to connect to many nodes
- ▶ Addresses of nodes are fetched during connection to the cluster - subset is specified pragmatically, full list resolved automatically
- ▶ Remember about that aspect during connecting to clusters
- ▶ Telepresence - might be usefull

# Concepts - Kafka

## Spark - streaming

- ▶ Take care of memory
- ▶ Java Kafka Client consumes direct memory - before deserialization
- ▶ Direct memory is not part of heap
- ▶ By default that part of memory is same as heap size

# Concepts - Custom source

## Spark - streaming

- ▶ No available in Python
- ▶ Implement simple interface
- ▶ onStart, onStop
- ▶ call store(element) for each element

# Concepts - Reliability

## Spark - streaming

- ▶ at-least-once, at-most-once, exactly-once
- ▶ Spark is observably exactly-once
- ▶ receivers can be reliable and unreliable - acknowledgements to the source


# Concepts - Transformation

## Spark - streaming

- ▶ Access underlying RDD
- ▶ Use-full when wanting to use static data with streams

# Concepts - Window operations

Spark - streaming



# Concepts - Window operations

## Spark - streaming

- ▶ tumbling
- ▶ sliding
- ▶ session


# Concepts - Window operations

## Spark - streaming

Tumbling Windows (5 mins)



Sliding Windows (10 mins, slide 5 mins)



Session Windows (gap duration 5 mins)





# Concepts - Window operations

Spark - streaming

Tumbling:

- ▶ Windows size

# Concepts - Window operations

Spark - streaming

Sliding:

- ▶ Windows size
- ▶ Interval length

# Concepts - Window operations

Spark - streaming

Session:

- ▶ Gap size

# Concepts - Joins

## Spark - streaming

- ▶ stream - stream - based on microbatches
- ▶ stream - data - using transformations

# Concepts - Outputs

## Spark - streaming

- ▶ files
- ▶ HDFS
- ▶ access each micro-batch RDD
- ▶ external libraries

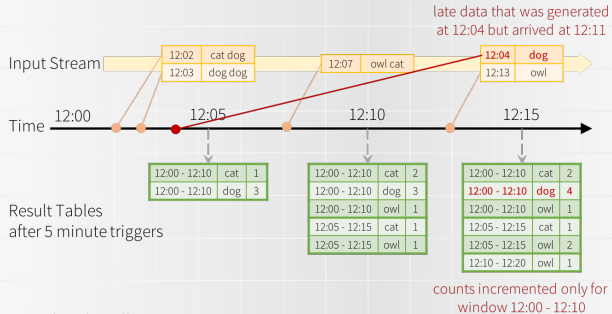
# Concepts - watermarking

## Spark - streaming

- ▶ how to handle late data
- ▶ spark is keeping track of current event time
- ▶ then it is used to manage internal states

# Concepts - watermarking

## Spark - streaming



Late data handling in  
Windowed Grouped Aggregation



# Next week

Spark - streaming

► Flink - streaming



# Large Scale Data Processing

## Lecture 4 – Spark - streaming

dr hab. inż. Tomasz Kajdanowicz, Piotr Bielak, Roman  
Bartusiak, Krzysztof Rajda

November 18, 2021