



Wrocław
University
of Science
and Technology

Large Scale Data Processing

Lecture 4 – Flink

dr hab. inż. Tomasz Kajdanowicz, Piotr Bielak, Roman
Bartusiak, Krzysztof Rajda

November 30, 2021





Overview


Flink



Flink

Actor programming

Flink



Actor programming

Flink

- ▶ Actors - can send message, modify hist state, spawn new actors
- ▶ Message queues
- ▶ Actor behaviour

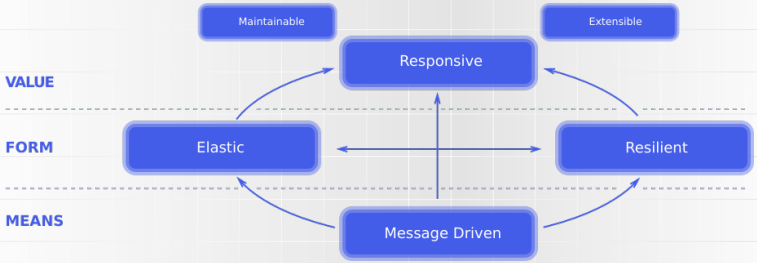
Actor programming

Flink

- ▶ Message - envelope with sender and receiver
- ▶ Offload response processing to third-actor
- ▶ Queues - communication decoupled from processing
- ▶ Only local state is mutable

Reactive manifesto

Flink



Akka platform

Flink



Akka platform

Flink

- ▶ Akka Actors
- ▶ Akka typed
- ▶ Akka streams
- ▶ Akka persistence
- ▶ Alpakka

Akka platform- Akka Actors

Flink

- ▶ core
- ▶ overhead - 300 bytes per actor
- ▶ actor system
- ▶ actor primitive
- ▶ communication patterns
 - ▶ tell
 - ▶ ask
 - ▶ forward
- ▶ coordination
- ▶ supervision

Akka platform- Akka typed

Flink

- ▶ new actor api
- ▶ typed communication
- ▶ typed behaviour
- ▶ typed state

Akka platform- Akka streams

Flink

- ▶ powerfull stream processing API
- ▶ processing graph DSL
- ▶ backpressure
- ▶ reactive streams
- ▶ materializes to actors

Akka platform- Akka cluster

Flink

- ▶ cluster formation
- ▶ cluster shutdown
- ▶ sharding
- ▶ singletons
- ▶ remote actors
- ▶ load balancing

Akka platform- Akka persistence

Flink

- ▶ event sourcing - next lecture
- ▶ persist actor state - by persisting events
- ▶ snapshot actor state - so the recover is not from all events
- ▶ recover actors
- ▶ latest state behaviour - persist only state not all events

Akka platform- Alpakka

Flink

- ▶ akka stream connectors
- ▶ Cassandra
- ▶ Kafka
- ▶ AWS

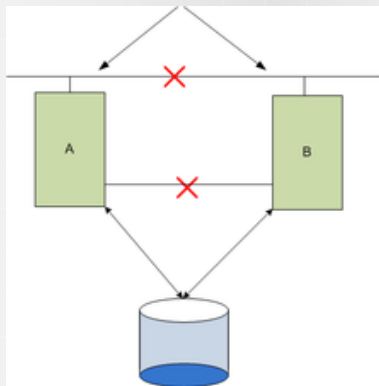
Akka cluster problems

Flink

- ▶ Split-brains
- ▶ Cluster formation
- ▶ Failure detection

Akka cluster - SB

Flink



Akka cluster - SBR

Flink

- ▶ Detect that some nodes are not responding
- ▶ Apply configured policy to handle split

Akka cluster - SBR - Keep majority

Flink

- ▶ cluster know it size
- ▶ on split, each part also knows size, and the size of new "cluster"
- ▶ down smaller part

Akka cluster - SBR - Static quorum

Flink

- ▶ configure quorum-size
- ▶ compare new "cluster" size to quorum
- ▶ down if size is smaller

Akka cluster - SBR - Keep oldest

Flink

- ▶ singletons are kept on oldest node
- ▶ on failure, if oldest node is not part of the new "cluster" - down

Akka cluster - SBR - Down all Flink

- ▶ kill the cluster

Akka cluster - SBR - Lease

Flink

- ▶ utilize external "lock"
- ▶ K8S
- ▶ time based
- ▶ lose lock if not updated in time

Partitioning

Flink

- ▶ Horizontal
- ▶ Vertical
- ▶ Functional


Partitioning - Horizontal

Flink

- ▶ Range
- ▶ Hash

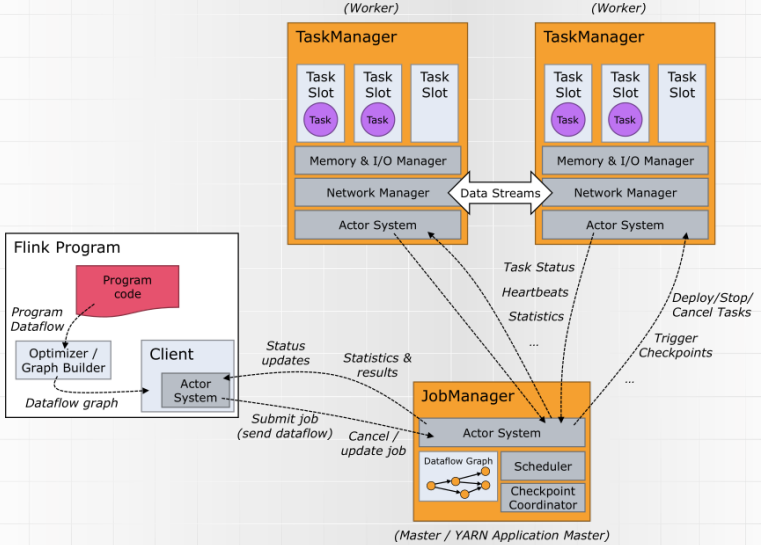
Partitioning

Flink



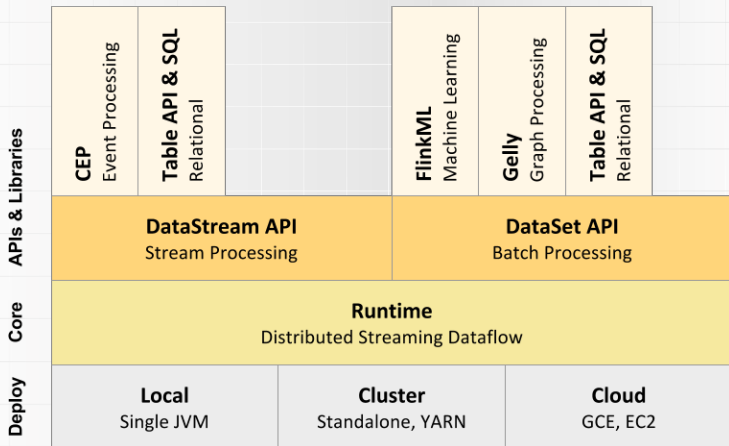
Architecture

Flink



Architecture - components

Flink



Concepts - Programs and DataFlows

Flink

- ▶ never ending flow of data
- ▶ transformations - operators
- ▶ "DAG"
- ▶ streaming dataflow - 1-n sources -> 1-n sinks

Concepts - Parallel DataFlows

Flink

- ▶ parallel by partitioning
- ▶ each operator - one or more subtask
- ▶ parallelism of stream = parallelism of upstream
- ▶ one-to-one - preserve partitioning
- ▶ redistributing - repartition

Concepts - Windows

Flink

- ▶ impossible to count all elements in a stream
- ▶ scope aggregation operations
- ▶ time drive
- ▶ data driven
- ▶ tumbling
- ▶ sliding
- ▶ session

Concepts - Time

Flink

- ▶ event time
- ▶ ingestion time
- ▶ processing time

Concepts - Statefull operations

Flink

- ▶ Simple processing - look only at current event
- ▶ remember information in between events - statefull operations
- ▶ partition by key
- ▶ only on "keyed" streams

Concepts - Checkpointing

Flink

- ▶ fault tolerance - reply + checkpointing
- ▶ checkpoint - might be expensive, do only in some intervals
- ▶ on failure - load checkpoint and replay events after the checkpoint

Concepts - Batch processing

Flink

- ▶ under the hood - streams
- ▶ no checkpoints
- ▶ in-memory state
- ▶ iterations

Concepts - Tasks and operators

Flink

- ▶ operator - for example map()
- ▶ task - chained operators, executed on single thread
- ▶ chaining - optimization
- ▶ configurable - programmable API

Concepts - Job managers, task managers, clients

Flink

- ▶ Job manager - master
- ▶ Task manager - worker
- ▶ Client - starts processing

Concepts - Task slots

Flink

- ▶ Task manager -Single JVM
- ▶ Task slot - part (or whole) part of task manager
- ▶ control on task isolation granularity
- ▶ task slot sharing
- ▶ needs at least as many task slots as max-parallelism

Concepts - State backend

Flink

- ▶ key/value store
- ▶ in-memory map
- ▶ RocksDB
- ▶ snapshotting as part of checkpoint

CEP

Flink

- ▶ Complex Event Processing
- ▶ on-top
- ▶ detect event patterns
- ▶ "regex" for event streams

Serialization

Flink

- ▶ by default internal efficient serialization for some primitive types, POJOs, case classes
- ▶ fallback - Kryo - schema in every event
- ▶ fallback - custom

Object reuse

Flink

- ▶ Flink - Java
- ▶ mutability
- ▶ user program can mutate objects, that is why Flink needs to copy them to be safe
- ▶ instruct Flink programmatically to reuse
- ▶ be immutable (copy where needed manually)

Next week

Flink

- ▶ Cassandra, CQRS, ES

Large Scale Data Processing

Lecture 4 - Flink

dr hab. inż. Tomasz Kajdanowicz, Piotr Bielak, Roman
Bartusiak, Krzysztof Rajda

November 30, 2021